



MODERN COMPUTER VISION

BY RAJEEV RATAN

Summary of Convolutional Neural Networks (CNNs)

Putting it all together

Conv Layers

$$(1 \times 0) + (0 \times 1) + (1 \times 0) + (1 \times 1) + (0 \times 0) + (0 \times -1) + (0 \times 0) + (1 \times 1) + (1 \times 0) = 2$$

1x0	0x1	1x0	0	1
1x1	0x0	0x-1	1	1
0x0	1x1	1x0	0	0
1	0	0	1	0
0	0	1	1	0

Input Image

*

0	1	0
1	0	-1
0	1	0

Filter or Kernel

=

2		

Output or Feature Map

- Convolution Operations occur when we Convolve our Filters with the input image by sliding it over our image
- This produces an output called a Feature Map
- Feature Maps are now the inputs to the next layer of our CNN

Stride, Padding and Kernel Size

$$\text{Feature Map Size} = n - f + 1 = m$$

$$\text{Feature Map Size} = 7 - 3 + 1 = 5$$

0	0	0	0	0	0	0
0	1	0	1	0	1	0
0	1	0	0	1	1	0
0	0	1	1	0	0	0
0	1	0	0	1	0	0
0	0	0	1	1	0	0
0	0	0	0	0	0	0

 $*$

0	1	0
1	0	-1
0	1	0

 $=$

2	1	-1	2	2
-1	1	3	2	1
2	1	1	1	2
1	1	1	0	2
2	0	2	3	1

7×7
 $n \times n$

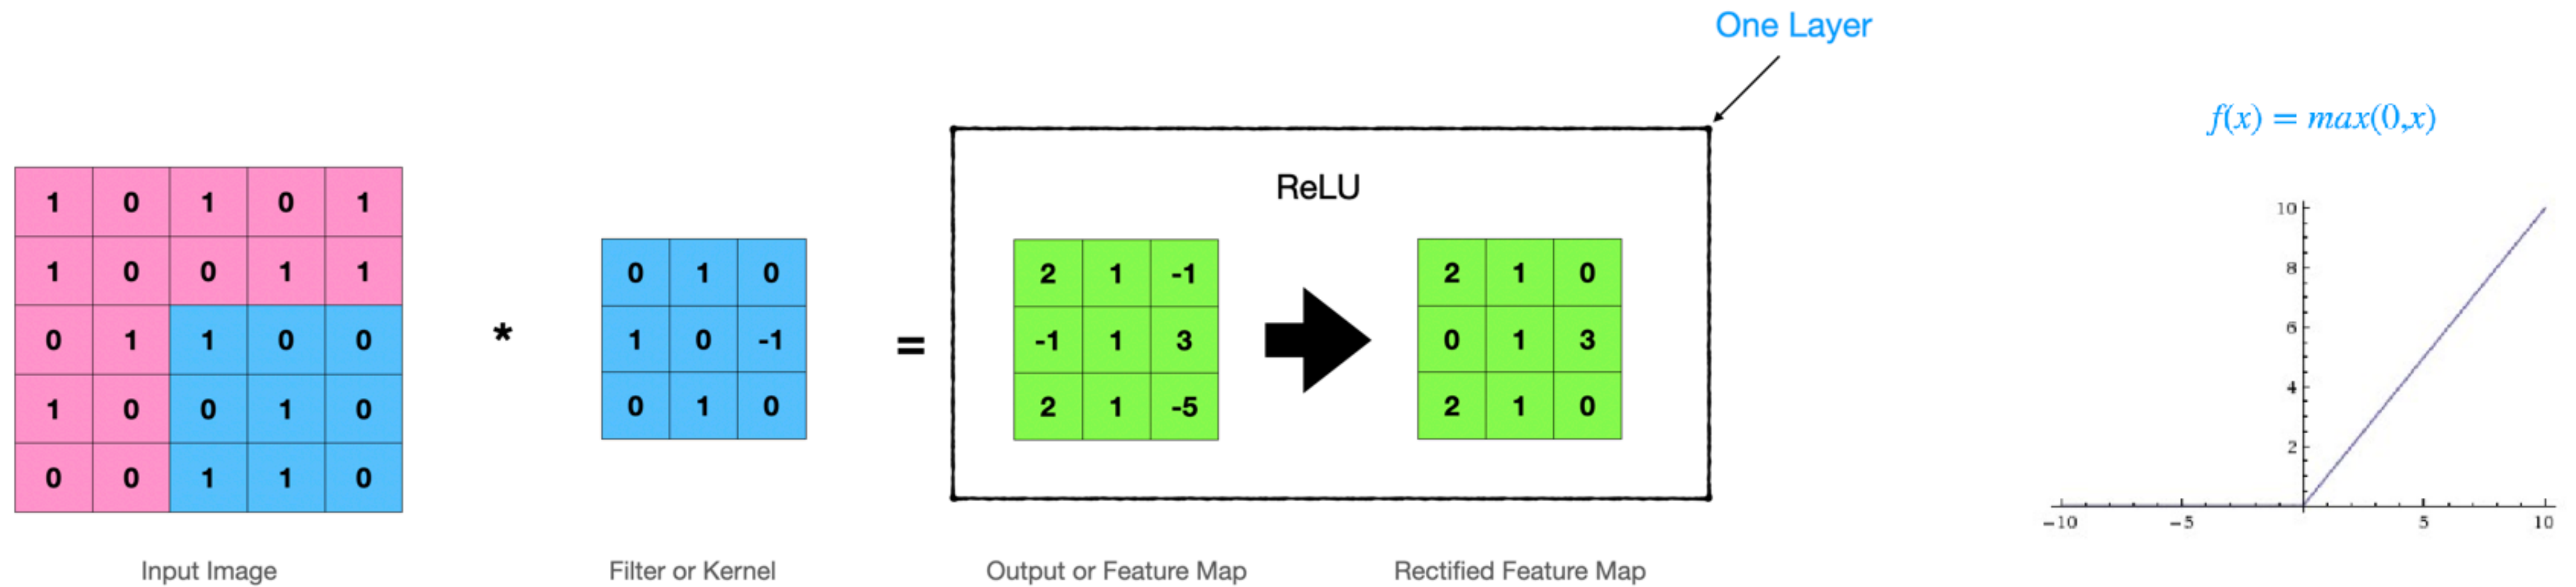
3×3
 $f \times f$

5×5
 $m \times m$

- We use Stride, Padding and Kernel size to control the output size of our Feature Map

$$(n \times n) * (f \times f) = \left(\frac{n + 2p - f}{s} + 1 \right) \times \left(\frac{n + 2p - f}{s} + 1 \right)$$

Activation Layer ReLU - Adds Non-Linearity to our Network



Max Pooling - Reduce Dimensionality

4	123	1	34
56	99	222	253
45	122	165	12
21	187	133	124

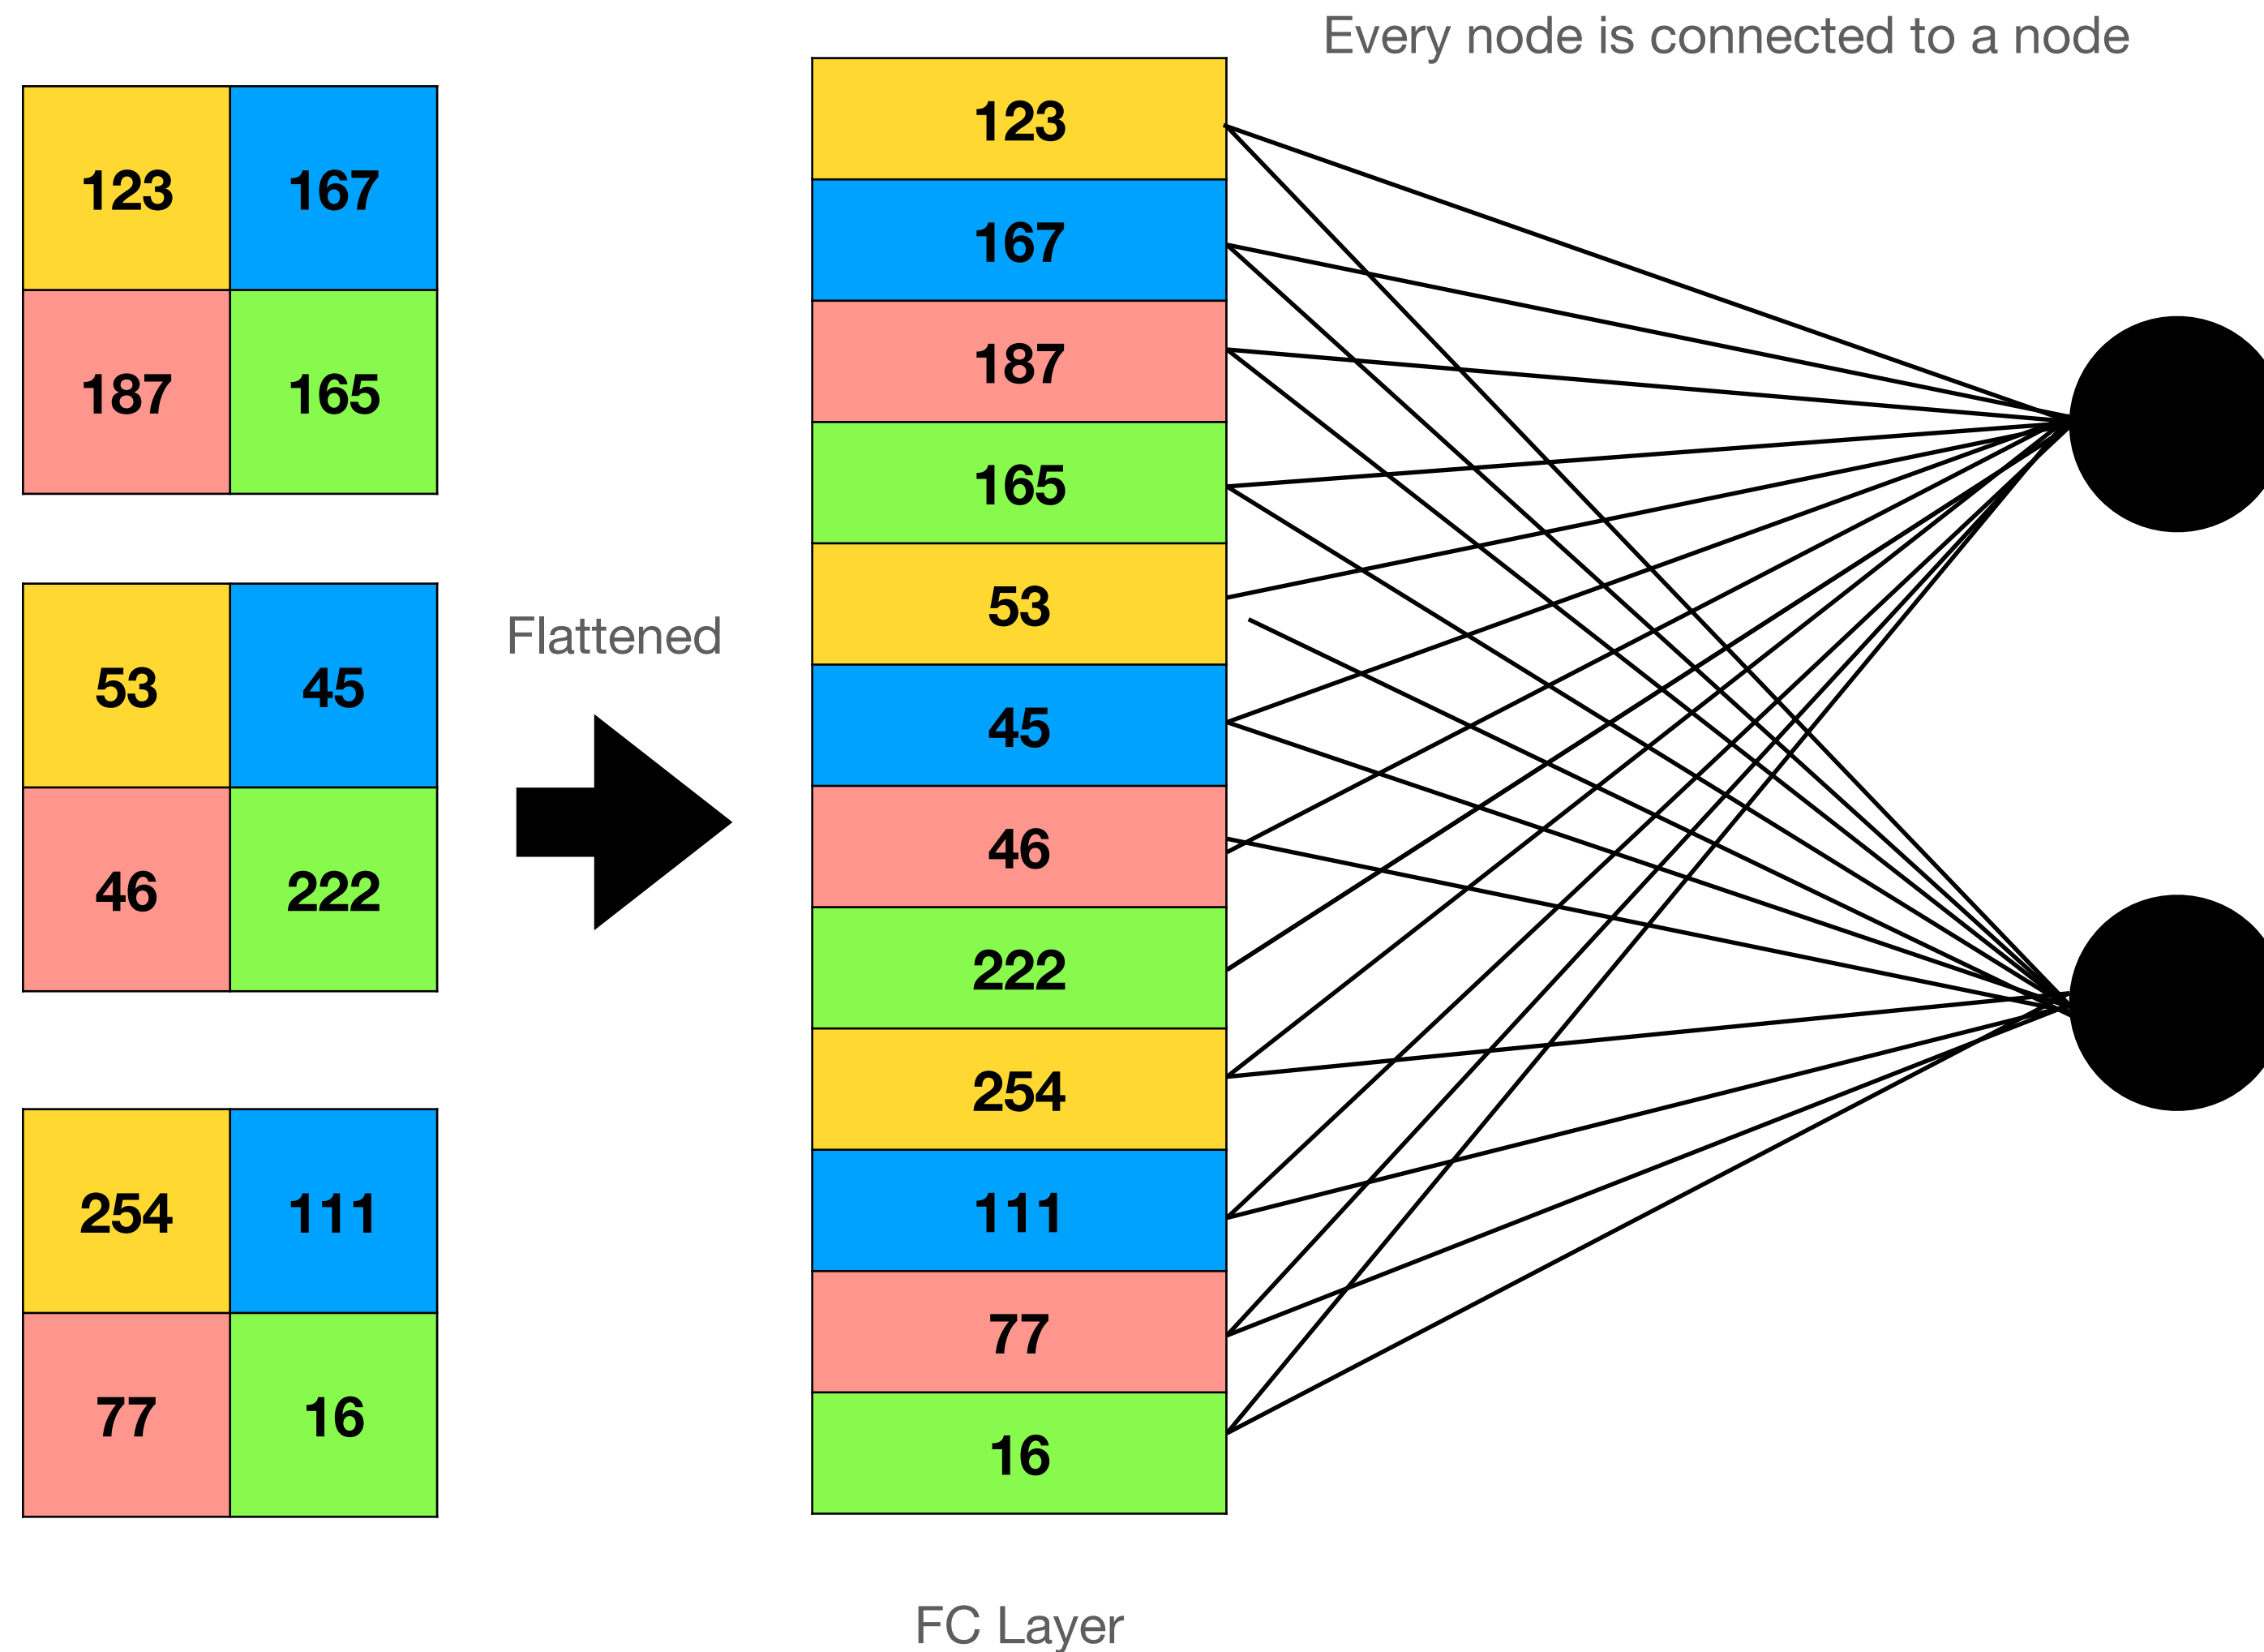
MaxPool Operation



Stride = 2
Kernel = 2x2

123	167
187	165

Fully Connected Layer - Max Pool Layer is Flattened



Softmax Layer

Logits Scores

2.0

1.0

0.1

Probabilities


0.7

0.2


0.1

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

Our Basic CNN



Parameters in our Basic CNN



Conv_1

$$((Height \times Width \times Depth) + bias) \times N_f$$

$$((3 \times 3 \times 1) + 1) \times 32 = 320$$

Conv_2

$$((Height \times Width \times Depth) + bias) \times N_f$$

$$((3 \times 3 \times 32) + 1) \times 64 = 18,494$$

No Trainable Parameters

- Max Pool
- Flatten
- ReLU

Fully Connected/Dense

$$(Length + bias) \times N_{nodes}$$

$$(9216 + 1) \times 128 = 1,179,776$$


Final Output (FC/Dense)

$$(Length + bias) \times N_{nodes}$$

$$(128 + 1) \times 10 = 1,290$$

Layer	Parameters
Conv_1 + ReLU	320
Conv_2 + ReLU	18494
Max Pool	0
Flatten	0
FC_1	1,179,776
FC_2 (Output)	1,290
Total	1,199,882

The Training Process



P	6	5	6	8	8	1
0	0	0.9	0.83	0.21	0.19	0.62
1	0.73	0.8	0.89	0.7	0.92	0.07
2	0.78	0.88	0.19	0.39	0.08	0.74
3	0.37	0.56	0.07	0.64	0.64	0.9
4	0.63	0.25	0.79	0.94	0.52	0.55
5	0.87	0.65	0.57	0.63	0.97	0.04
6	0.67	0.05	0.45	0.51	0.87	0.51
7	0.71	0.66	0.13	0.59	0.86	0.89
8	0.51	0.88	0.59	0.01	0.37	0.63
9	0.24	0.52	0.79	0.15	0.63	0.78

Overview on Training

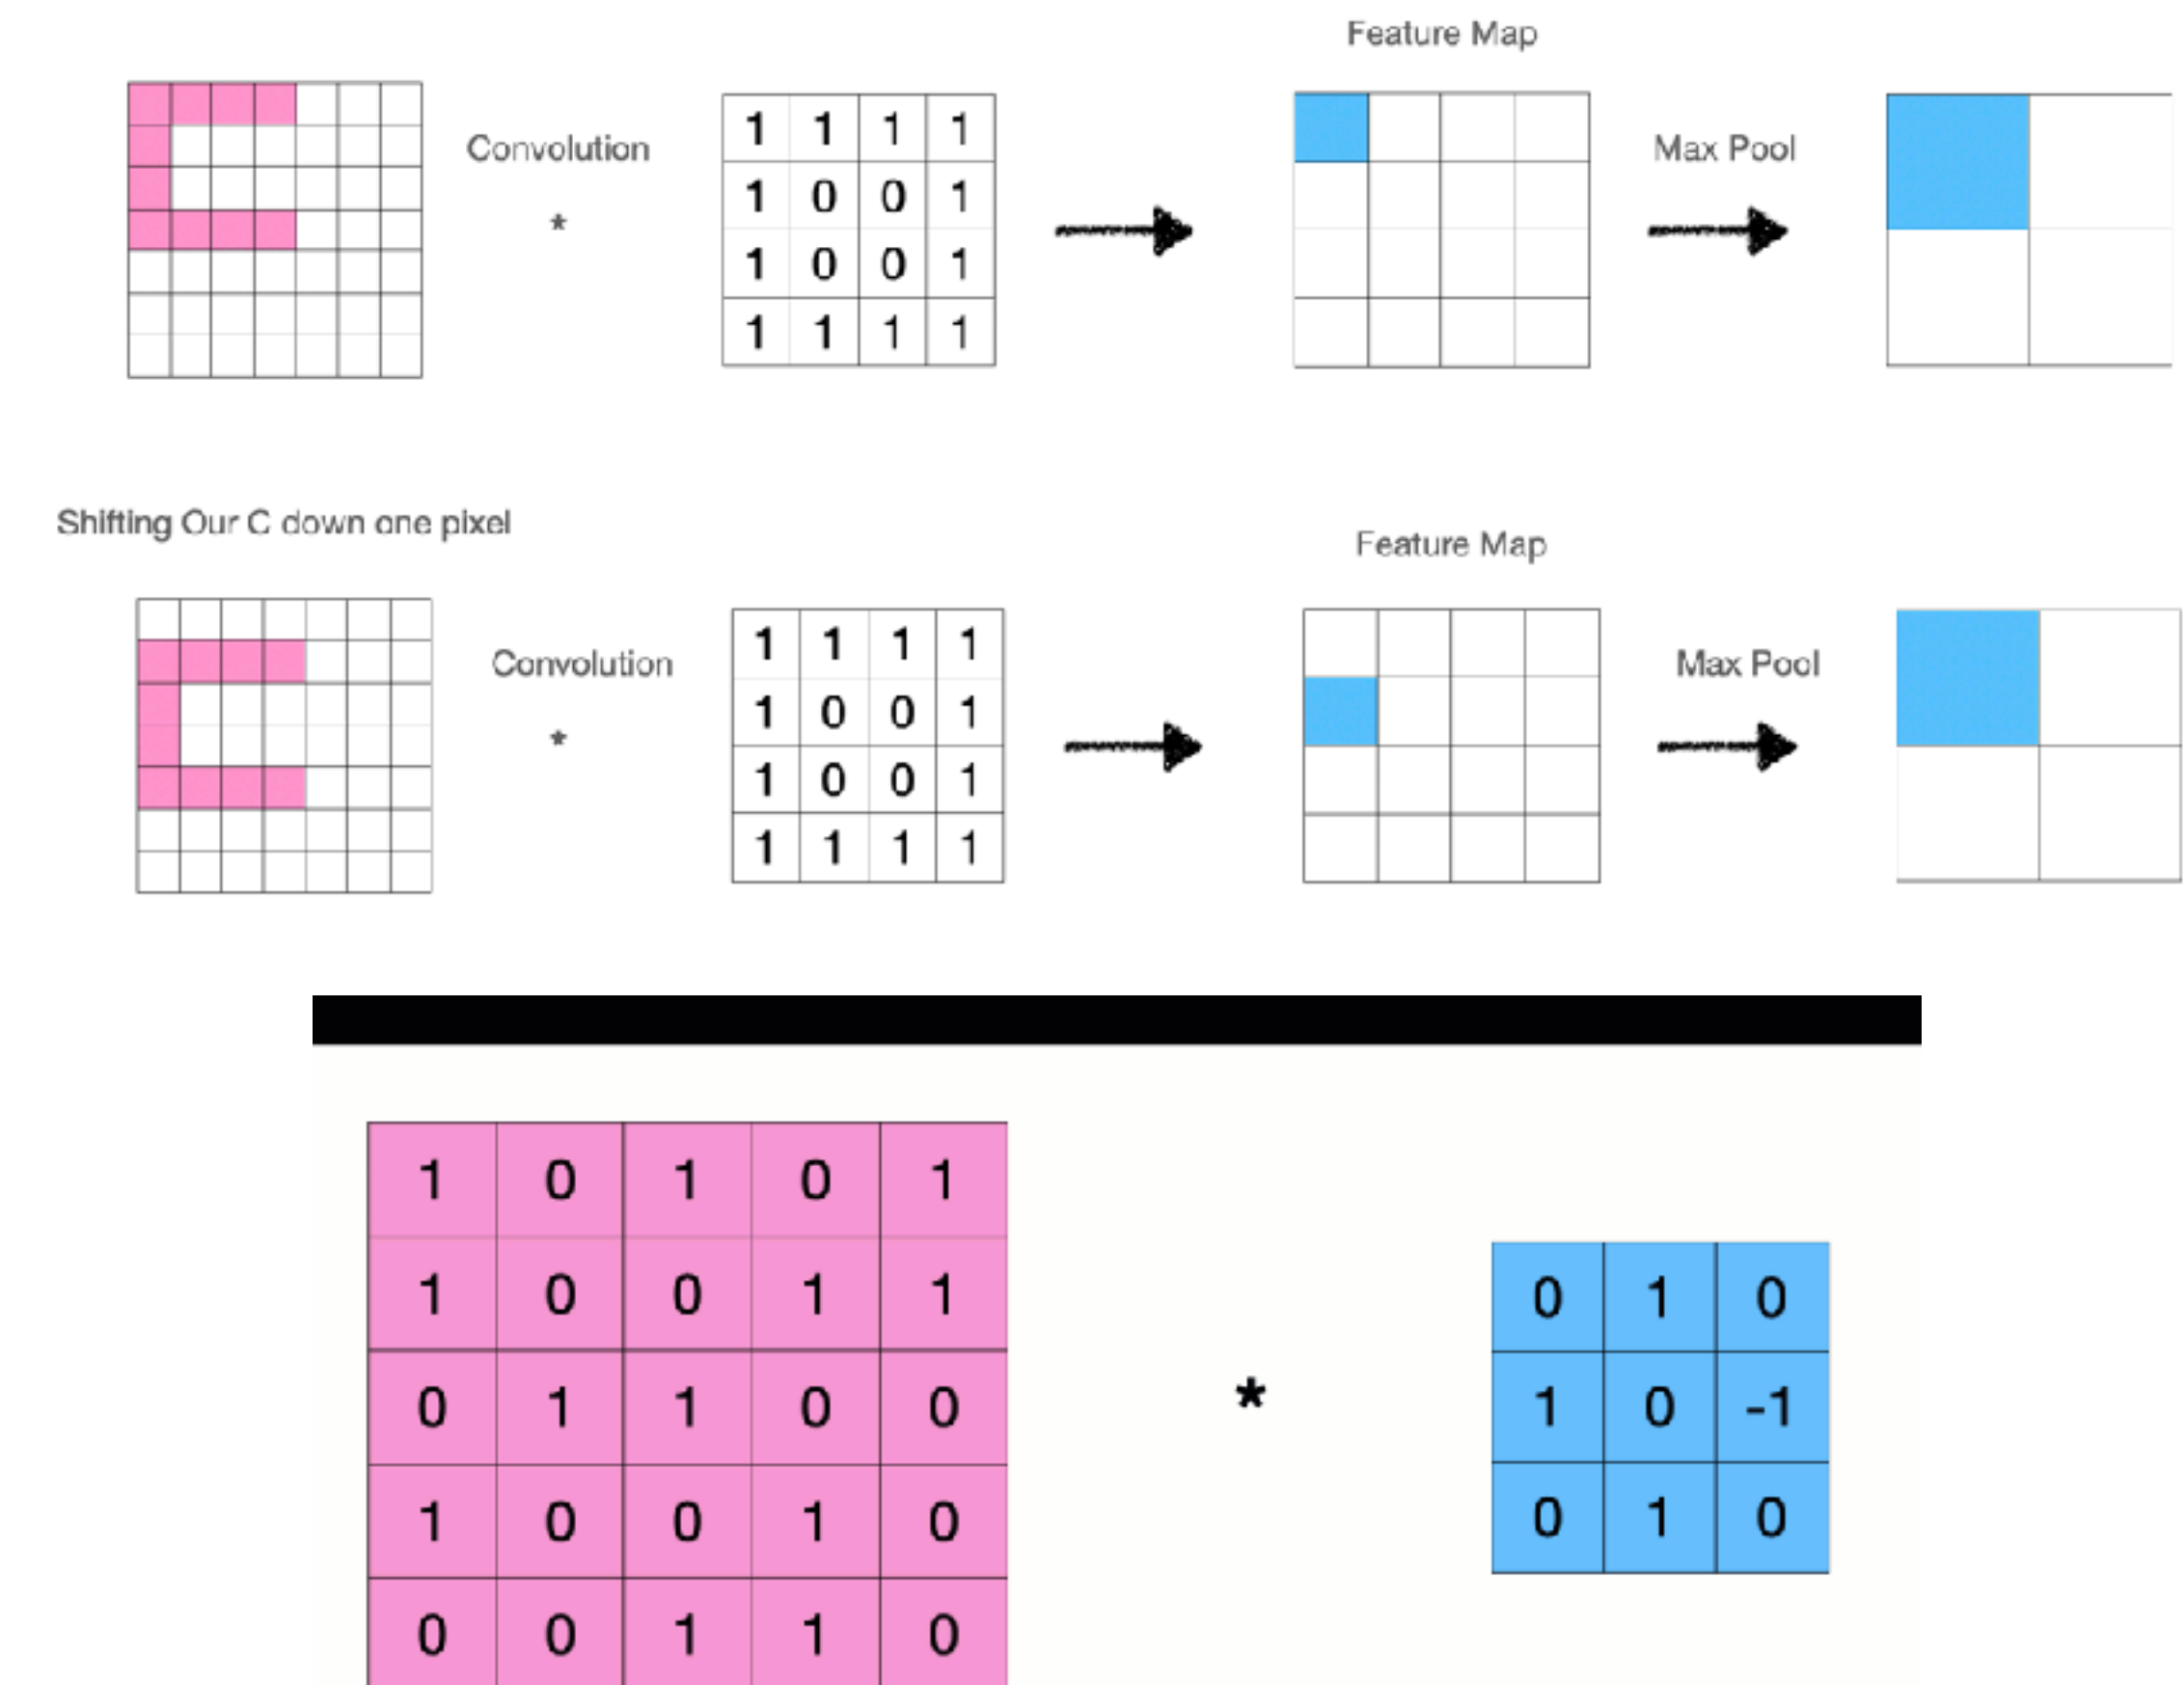
- CNN Model is **designed and defined**
- Weights are **initialised with random values**
- Batches of Images (typically 8 to 256) are **forward propagated** through our CNN Model
- Using **Back Propagation** with **Mini-Batch Gradient Descent** we update the individual weights (right to left)
- Using we update all our weights so that we have a lower loss
- We the entire dataset of images is forward propagated, we've completed an **Epoch**
- We train for **5 to 50** Epochs and stop when Loss stops decreasing

Batches, Mini-Batches, Iterations & Epochs

- We can feed images one at a time, however using mini-batches is better


Advantages of Convolution Neural Networks

- **Invariance** - Remember our Max Pool Example
- **Parameter sharing** - where a single filter can be used all parts of an image
- **Sparsity of connections** - As we saw, fully connected layers in a typical Neural Network result in a weight matrix with large number of parameters.



Convolution Neural Networks Assumptions

- Low-level features are local
- Features are translational invariant
- High-level features are made up of low-level features





MODERN COMPUTER VISION

BY RAJEEV RATAN

Next...

History of Deep Learning and AI